



INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

# Introduction to the Course

**Rob Reider**

Adjunct Professor, NYU-Courant  
Consultant, Quantopian




# Example of Time Series: Google Trends







# Example of Time Series: Climate Data



# Example of Time Series: Quarterly Earnings Data



# Example of Multiple Series: Natural Gas and Heating Oil





# Goals of Course

- Learn about time series models
- Fit data to a times series model
- Use the models to make forecasts of the future
- Learn how to use the relevant statistical packages in Python
- Provide concrete examples of how these models are used



# Some Useful Pandas Tools

- Changing an index to datetime

```
df.index = pd.to_datetime(df.index)
```

- Plotting data

```
df.plot()
```

- Slicing data

```
df['2012']
```



# Some Useful Pandas Tools

- Join two DataFrames

```
df1.join(df2)
```

- Resample data (e.g. from daily to weekly)

```
df = df.resample(rule='W', how='last')
```



# More pandas Functions

- Computing percent changes and differences of a time series

```
df['col'].pct_change()  
df['col'].diff()
```

- pandas correlation method of Series

```
df['ABC'].corr(df['XYZ'])
```

- pandas autocorrelation

```
df['ABC'].autocorr()
```



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

**Let's practice!**



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

# Correlation of Two Time Series


**Rob Reider**

Adjunct Professor, NYU-Courant  
Consultant, Quantopian



# Correlation of Two Time Series

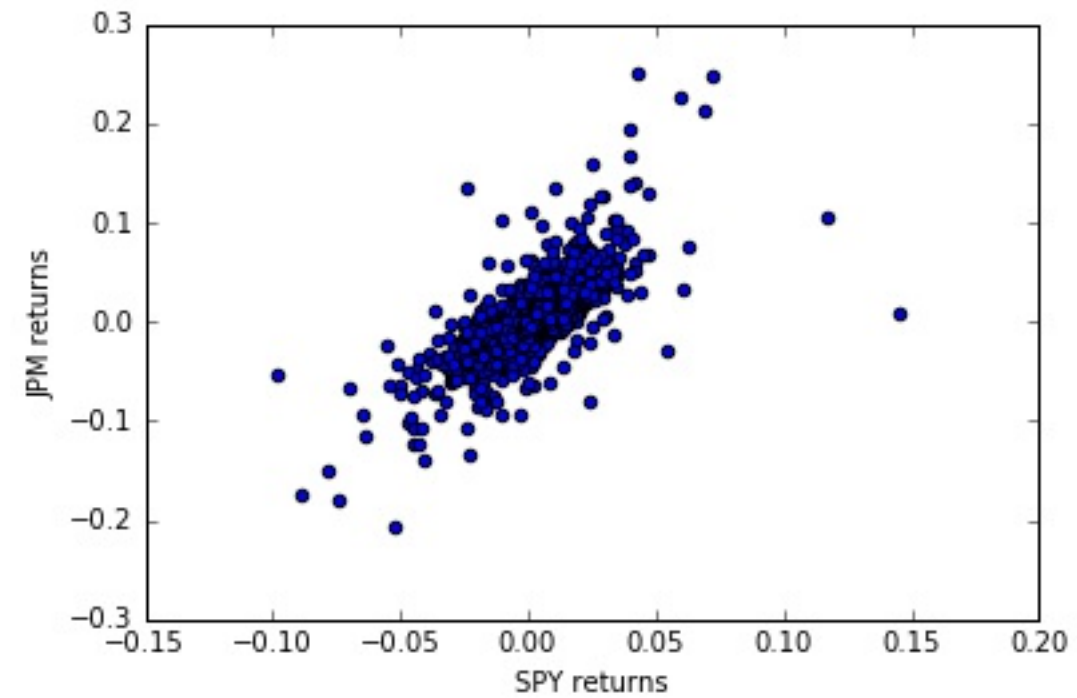
- Plot of S&P500 and JPMorgan stock





# Correlation of Two Time Series


- Scatter plot of S&P500 and JP Morgan returns






# More Scatter Plots


- Correlation = 0.9




- Correlation = 0.4



- Correlation = -0.9




- Correlation = 1.0





# Common Mistake: Correlation of Two Trending Series

- Dow Jones Industrial Average and UFO Sightings ([www.nuforc.org](http://www.nuforc.org))



- Correlation of levels: 0.94
- Correlation of percent changes:  $\approx 0$



# Example: Correlation of Large Cap and Small Cap Stocks

- Start with stock prices of SPX (large cap) and R2000 (small cap)
- First step: Compute percentage changes of both series

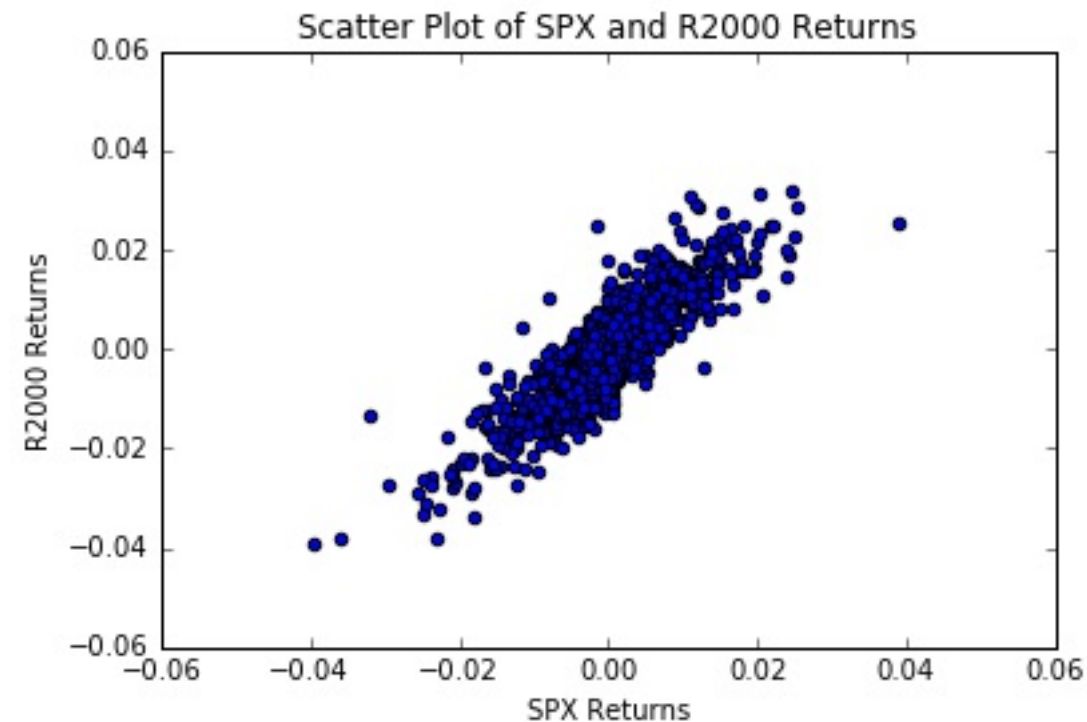
```
df['SPX_Ret'] = df['SPX_Prices'].pct_change()  
df['R2000_Ret'] = df['R2000_Prices'].pct_change()
```



# Example: Correlation of Large Cap and Small Cap Stocks

- Visualize correlation with scatter plot

```
plt.scatter(df['SPX_Ret'], df['R2000_Ret'])  
plt.show()
```





# Example: Correlation of Large Cap and Small Cap Stocks

- Use pandas correlation method for Series

```
correlation = df['SPX_Ret'].corr(df['R2000_Ret'])  
print("Correlation is: ", correlation)
```

```
Correlation is: 0.868
```



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

**Let's practice!**



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

# Simple Linear Regressions

**Rob Reider**

Adjunct Professor, NYU-Courant  
Consultant, Quantopian



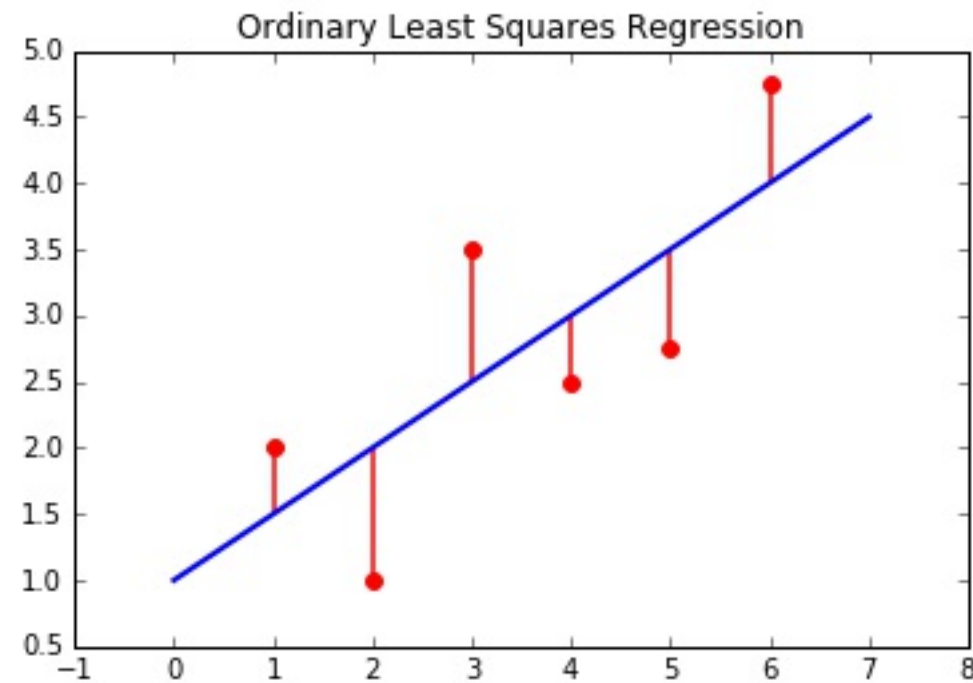
# What is a Regression?

- **Simple linear regression:**

$$y_t = \alpha + \beta x_t + \epsilon_t$$

# What is a Regression?

- **Ordinary Least Squares (OLS)**



# Python Packages to Perform Regressions

- In statsmodels:

```
import statsmodels.api as sm  
sm.OLS(y, x).fit()
```

- In numpy:

```
np.polyfit(x, y, deg=1)
```

- In pandas:

```
pd.ols(y, x)
```

- In scipy:

```
from scipy import stats  
stats.linregress(x, y)
```

- Beware that the order of x and y is not consistent across packages

# Example: Regresssion of Small Cap Returns on Large Cap

- Import the statsmodels module

```
import statsmodels.api as sm
```

- As before, compute percentage changes in both series

```
df['SPX_Ret'] = df['SPX_Prices'].pct_change()  
df['R2000_Ret'] = df['R2000_Prices'].pct_change()
```

- Add a constant to the DataFrame for the regression intercept

```
df = sm.add_constant(df)
```



# Regression Example (continued)

- Notice that the first row of returns is NaN

| Date       | SPX_Price   | R2000_Price | SPX_Ret   | R2000_Ret |
|------------|-------------|-------------|-----------|-----------|
| 2012-11-01 | 1427.589966 | 827.849976  | NaN       | NaN       |
| 2012-11-02 | 1414.199951 | 814.369995  | -0.009379 | -0.016283 |

- Delete the row of NaN

```
df = df.dropna()
```

- Run the regression

```
results = sm.OLS(df['R2000_Ret'], df[['const', 'SPX_Ret']]).fit()  
print(results.summary())
```



# Regressionssion Example (continued)

- Regression output

```
OLS Regression Results
=====
Dep. Variable:      R2000_Ret      R-squared:          0.753
Model:              OLS           Adj. R-squared:      0.753
Method:             Least Squares  F-statistic:        3829.
Date:               Fri, 26 Jan 2018 Prob (F-statistic):  0.00
Time:               13:29:55       Log-Likelihood:     4882.4
No. Observations:   1257          AIC:                -9761.
Df Residuals:       1255          BIC:                -9751.
Df Model:           1
Covariance Type:    nonrobust
=====
                    coef    std err          t      P>|t|      [95.0% Conf. Int.]
-----
const      -4.964e-05    0.000      -0.353    0.724      -0.000      0.000
SPX_Ret      1.1412    0.018     61.877    0.000       1.105      1.177
=====
Omnibus:            61.950   Durbin-Watson:       1.991
Prob(Omnibus):      0.000   Jarque-Bera (JB):    148.100
Skew:               0.266   Prob(JB):            6.93e-33
Kurtosis:           4.595   Cond. No.            131.
=====
```

- Intercept in results.params[0]
- Slope in results.params[1]

# Regressionssion Example (continued)

- Regression output

```
OLS Regression Results
=====
Dep. Variable:      R2000_Ret      R-squared:      0.753
Model:              OLS           Adj. R-squared:  0.753
Method:             Least Squares  F-statistic:    3829.
Date:               Fri, 26 Jan 2018  Prob (F-statistic): 0.00
Time:               13:29:55       Log-Likelihood:  4882.4
No. Observations:   1257          AIC:             -9761.
Df Residuals:       1255          BIC:             -9751.
Df Model:           1
Covariance Type:    nonrobust
=====
                    coef    std err          t      P>|t|      [95.0% Conf. Int.]
-----
const      -4.964e-05    0.000      -0.353    0.724      -0.000    0.000
SPX_Ret      1.1412    0.018     61.877    0.000      1.105    1.177
=====
Omnibus:            61.950   Durbin-Watson:      1.991
Prob(Omnibus):      0.000   Jarque-Bera (JB):    148.100
Skew:               0.266   Prob(JB):            6.93e-33
Kurtosis:           4.595   Cond. No.            131.
=====
```



# Relationship Between R-Squared and Correlation

- $[\text{corr}(x, y)]^2 = R^2$  (or R-squared)
- $\text{sign}(\text{corr}) = \text{sign}(\text{regression slope})$
- In last example:
  - R-Squared = 0.753
  - Slope is positive
  - correlation =  $+\sqrt{0.753} = 0.868$



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

**Let's practice!**



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

# Autocorrelation

**Rob Reider**

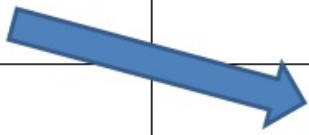
Adjunct Professor, NYU-Courant  
Consultant, Quantopian



# What is Autocorrelation?

- Correlation of a time series with a lagged copy of itself

| Series | Lagged Series |
|--------|---------------|
| 5      |               |
| 10     | 5             |
| 15     | 10            |
| 20     | 15            |
| 25     | 20            |
| ⋮      | ⋮             |

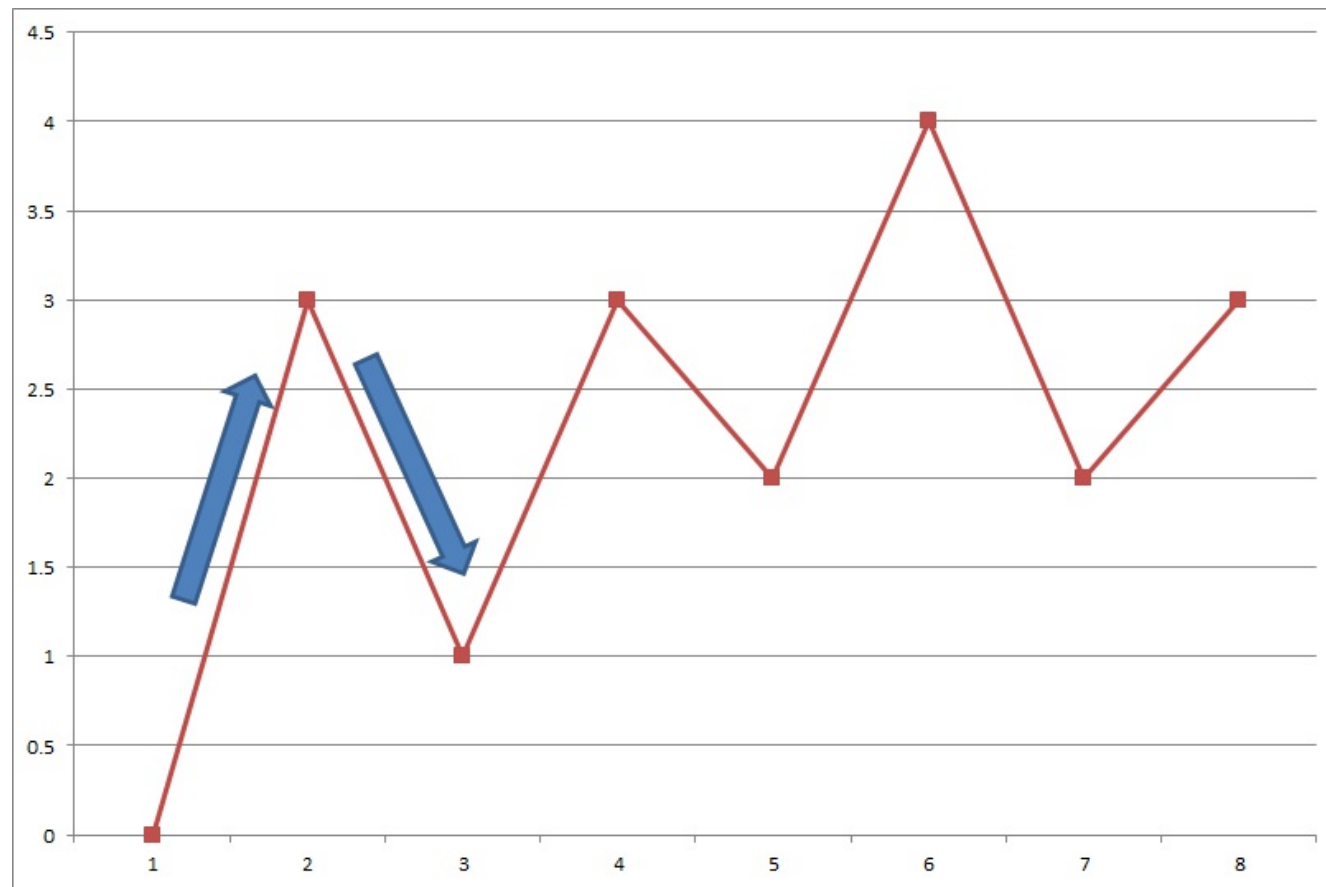


- Lag-one autocorrelation
- Also called **serial correlation**



# Interpretation of Autocorrelation

- **Mean Reversion** - Negative autocorrelation

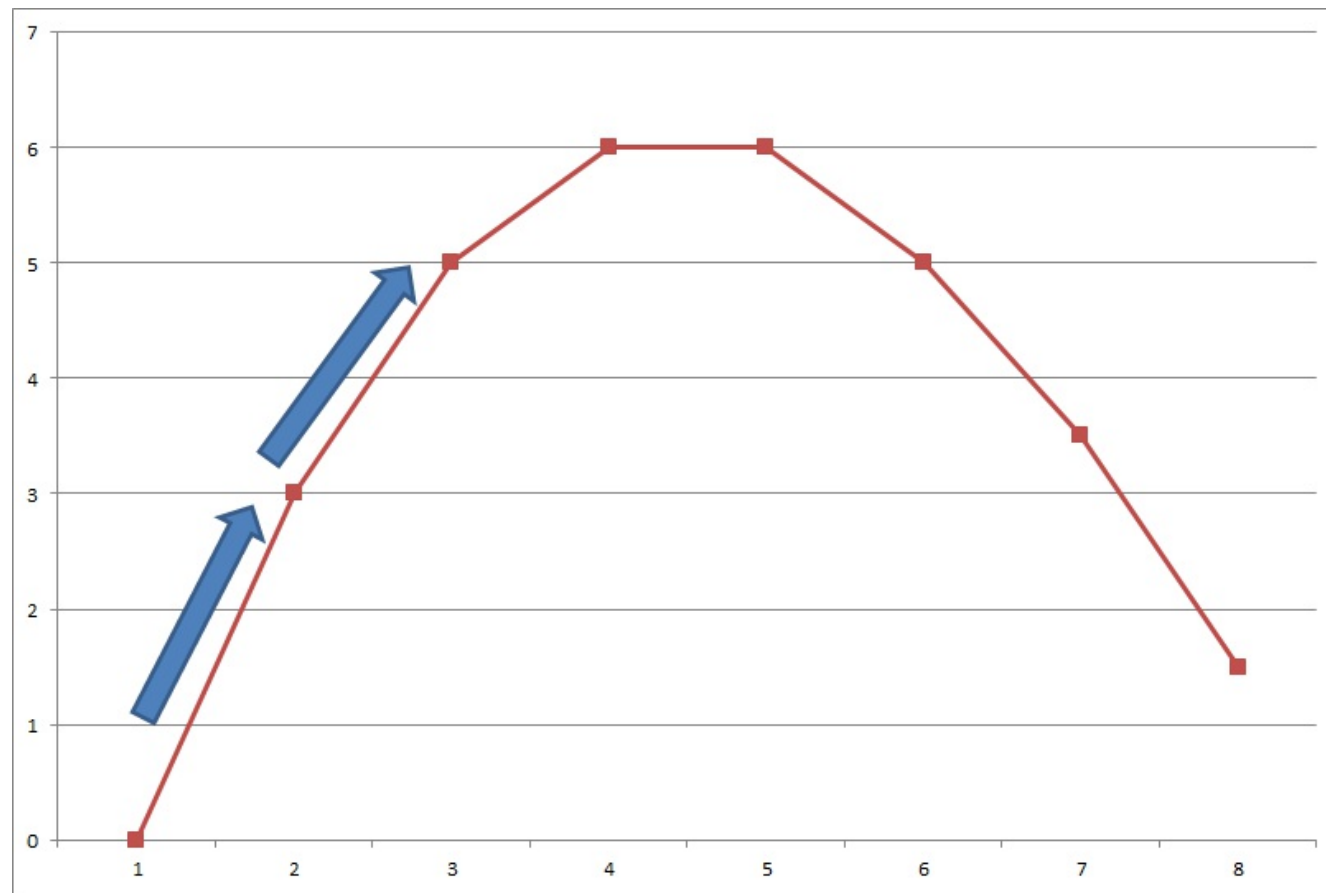






# Interpretation of Autocorrelation

- **Momentum, or Trend Following** - Positive autocorrelation





# Traders Use Autocorrelation to Make Money

- Individual stocks
  - Historically have negative autocorrelation
  - Measured over short horizons (days)
  - Trading strategy: Buy losers and sell winners
- Commodities and currencies
  - Historically have positive autocorrelation
  - Measured over longer horizons (months)
  - Trading strategy: Buy winners and sell losers

# Example of Positive Autocorrelation: Exchange Rates

- Start with daily data of ¥/\$ exchange rates in DataFrame `df` from

[FRED](#)

- Convert index to datetime

```
df.index = pd.to_datetime(df.index)
```

- Downsample from daily to monthly data

```
df = df.resample(rule='M', how='last')
```

- Compute returns from prices

```
df['Return'] = df['Price'].pct_change()
```

- Compute autocorrelation

```
autocorrelation = df['Return'].autocorr()  
print("The autocorrelation is: ", autocorrelation)  
The autocorrelation is: 0.0567
```



## INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

**Let's practice!**